

Ch 04: Traps and System Calls

xv6 a simple Unix-like teaching Operating System

Derek Harter

Department of Computer Science
East Texas A&M University

Summer 2026



EAST TEXAS A&M

- 1 Traps and System Calls
- 2 Traps from User Space
- 3 Traps from Kernel Space



Ch 04: Traps and System Calls

xv6 a simple Unix-like teaching Operating System

Derek Harter

Department of Computer Science
East Texas A&M University

Summer 2026



EAST TEXAS A&M

Traps and System Calls (1 of 3)

There are 3 kinds of event which cause the CPU to set aside ordinary execution of instructions and force a transfer of control to special kernel code that handles the event:

- ① **system call**: a user program executes the `ecall` (RISC-V) instruction to ask the kernel to do something for it
- ② **exception**: an instruction (user or kernel mode) does something illegal, such as attempt to load an invalid virtual address
- ③ device **interrupt**: when a device signals that it needs attention



Introduction Traps and System Calls (2 of 3)

- We will use **trap** as a generic term for all 3 situations
- In all cases, whatever code was executing at the time of a trap will later need to resume
 - Trap call needs to be **transparent**, interrupted code does not know it happened
- A trap forces transfer of control into the kernel
 - kernel saves registers and other state (PC) so can resume execution later
 - kernel executes appropriate handler code (system call, exception handler or device driver)
 - kernel restores the saved state and returns from trap once handled, original code (might) resume where it left off



Introduction Traps and System Calls (3 of 3)

xv6 trap handling proceeds in four? stages:

- ① hardware actions taken by the RISC-V cpu
- ② some assembly instructions that prepare the way for kernel C code trap handlers
- ③ a C function that decides what to do with the trap
- ④ finally the system call or device-driver service routine
 - implicitly state is saved in 1,2, and restored at end of 4

Kernel code that processes a trap is called a **handler**

The first handler instructions are usually written in assembler and are sometimes called a **vector**



RISC-V Trap Machinery

RISC-V CPU has a set of hardware control registers that the kernel writes to tell the CPU how to handle traps, and that the kernel can read to find out about a trap that has occurred

- **stvec**: kernel writes address of its trap handler code here, RISC-V CPU jumps to the address in `stvec` to handle a trap
- **sepc**: when a trap occurs, RISC-V saves the program counter here
 - `sret` (return from trap) copies `sepc` to `pc`, kernel can write `sepc` to control where `sret` goes
- **scause**: RISC-V puts a number here that describes the reason for the trap
- **sscratch**: kernel trap handler code uses `sscratch` to help avoid it overwriting user registers before saving them
- **sstatus**: The SIE bit in `sstatus` controls whether device interrupts are enabled. If kernel clears SIE RISC-V will defer device interrupts until it is set again.
 - The SPP bit indicates whether a trap came from user mode or supervisor mode, and controls what mode `sret` returns to.



RISC-V Hardware Actions

When it forces a trap, the RISC-V hardware does the following

- 1 If the trap is a device interrupt, and the `sstatus` SIE bit is clear, don't do any following
- 2 Otherwise, disable interrupts by clearing the SIE bit in `sstatus`
- 3 Copy the `pc` to `sepc`
- 4 Save the current mode (user or supervisor) in the SPP bit in `sstatus`
- 5 Set `scause` to a number indicating the trap's cause
- 6 Set the mode to supervisor
- 7 Copy `stvec` to the `pc`
- 8 Start executing at the new `pc`

These actions are hardwired into the CPU, but note RISC-V doesn't do the following, which are sometimes done on other hardware architectures:

- Does not switch to the kernel page table (so will still be using user/process page table)
- Doesn't switch to a stack in the kernel (so still using user stack initially)
- Doesn't save any registers, other than the `pc`



Summary

- 3 kinds of events cause a **trap** into supervisor trap handler: **system call** (`ecall`), **exception**, device **interrupt**
- A trap forces hardware transfer of control into the kernel, supervisor mode trap handler
- RISC-V uses registers like `stvec` and `sepc` to allow software to enable and handle hardware generated traps, and switch from user mode to kernel mode safely
- RISC-V is hardwired to save the `pc`, communicate the cause of the interrupt, and load the trap vector specified by the OS into the `pc` to execute and return from traps
 - Allows the kernel OS software to setup how to handle traps in supervisor mode



Ch 04: Traps from User Space

xv6 a simple Unix-like teaching Operating System

Derek Harter

Department of Computer Science
East Texas A&M University

Summer 2026



EAST TEXAS A&M

Traps from User Space

xv6 handles traps differently depending on whether the trap occurs while executing in kernel or in user code

A trap may occur for any of the 3 reasons from user space, program makes a system call (`ecall` instruction), does something illegal, or a device interrupts.

- `stvec` is set to jump to `uservec` in the trampoline page (all address spaces map here)
 - saves all registers in `TRAPFRAME`
 - switches to the kernel page table
- calls into `usertrap()`
 - determines and handles system call, device interrupt or exception
- `prepare_return()` and ultimately `userret` are called in trampoline, restores registers and returns to user mode




Figure 1: Outline of how a trap from user code is handled (Cox et al., 2026, pg.45)



Calling System Calls (1 of 2)

- User programs call library functions in order to make system calls, for example

```
write(2, "$ ", 2); // write stdout fd=2, the string, length 2 chars
```

- You can find the library function in `user/usys.S`:

```
write:  
    li a7, SYS_write  
    ecall  
    ret
```

- The 3 parameters to write are passed in registers a0, a1 and a2
 - a7 holds the system call number, it is set by the call to write as shown
- The `ecall` instruction traps from user space causing `uservec usertrap()` then `syscall()` to execute



Calling System Calls (2 of 2)

- `syscall()` retrieves the system call number from `a7` and uses to index into `syscalls`
 - for example `SYS_write` indexes and calls the `sys_write()` function in the kernel
- When `sys_write()` returns, `syscall` records its return value in `p->trapframe->a0`
 - This will be what the original user-space call to `write()` will be returned
 - By convention 0 is success, and non zero is failure? (I disagree with book a bit here).



System Call Arguments

- System call arguments start out in user registers `a0 a1 ...` and are moved to the trap frame by the kernel trap code
- kernel functions `argint()` `argaddr()` and `argfd()` can be used to retrieve the *n*'th system call argument from the trap frame as an integer, pointer, or a file descriptor
- Sometimes a pointer is passed, will be a pointer to user-space address space
 - Has to be checked, could be buggy or malicious
 - Also the kernel page table mappings are not the same as the user page table, so kernel cannot use ordinary instructions to load or store from user-supplied addresses
- kernel implements functions that safely transfer data to and from user-supplied addresses
 - `fetchstr()` retrieve string file-name arguments
 - `copyinsstr()` copy in `str` copies up to `max` bytes to `dst` from virtual address `srvca` in the user page table `pagetable`



Summary

- xv6 handles traps from user space by entering the `uservec` in the trampoline page
 - This is mapped to same virtual space in all user page tables and the kernel page table
 - basically saves the user process state from where it was interrupted
- The `uservec` trap vector calls the `usertrap()` routine to handle the actual trap
 - This does some things like reenables interrupts and switch to the kernel stack and kernel page table
 - handles system call, exception or interrupt as needed
 - calls `prepare_return()` to setup for return, including switching back to the process page table and stack, and ensuring that the `epc` register has address to return back pc to
- When invoking an xv6 system call
 - the `ecall` instruction performs a software trap to change from user mode to supervisor mode
 - system call parameters are placed in `a0`, `a1`, `a2`, ... registers
 - `a7` register is loaded with the index of the system call to perform
 - The `usertrap()` function calls `syscall()` to execute a system call, which looks up `a7` in the trapframe to determine which system call was invoked
 - The return result from a system call is returned in `a0`, which is first placed in the saved `trap` frames `a0` location so it is restored to actual register `a0` on return to user space



Ch 04: Traps from Kernel Space

xv6 a simple Unix-like teaching Operating System

Derek Harter

Department of Computer Science
East Texas A&M University

Summer 2026



EAST TEXAS A&M

Traps from Kernel Space

- Traps while in kernel mode are handled in a different way than traps from user space
- `stvec` will point to `kernelvec`
 - Can rely on only executing when already in kernel mode, so don't need to switch kernel page table nor kernel stack
 - `kernelvec` saves registers needed on the kernel stack (rather than `trapframe`) which makes sense because
 - also this allows for a trap to happen when already in a trap, since pushing on stack
 - also for multicore, this allows to handle traps that occur on different cpus
- Only device interrupts are handled as traps from kernel traps.
 - Obviously we don't need system calls, that is a mechanism to switch from user space to kernel space
 - Exceptions here will just cause a `kernel panic()`, they cannot be recovered from as an exception from user program might sometimes be able to



Summary



Cox, R., Kaashoek, F., & Morris, R. (2026). *Xv6: A simple, unix-like teaching operating system*. Retrieved September 2, 2025, from <https://github.com/mit-pdos/xv6-riscv-book>

